



## Part B - Lab Questions

In [119... `"""Q1.(a) Create a 1D NumPy array with values 1 to 20.  
(b) Extract all prime numbers from it.  
(c) Compute the mean and variance of the extracted primes."""`

```
# (a) Create a 1D NumPy array with values 1 to 20
import numpy as np
ar = np.arange(1,21)
print("Array is:",ar)

# (b) Extract all prime numbers
def is_prime(n):
    if n < 2:
        return False
    for i in range(2, int(np.sqrt(n)) + 1):
        if n % i == 0:
            return False
    return True

primes = np.array([x for x in ar if is_prime(x)])
print("Prime numbers:", primes)

# (c) Compute mean and variance of primes
mean_primes = np.mean(primes)
var_primes = np.var(primes)
print("Mean of primes:", mean_primes)
print("Variance of primes:", var_primes)
```

Array is: [ 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20]  
Prime numbers: [ 2 3 5 7 11 13 17 19]  
Mean of primes: 9.625  
Variance of primes: 35.734375

In [120... `"""Q2.(a) Create a 4x4 NumPy array with numbers 1 to 16.  
(b) Extract the 2x2 bottom-left sub-matrix.  
(c) Compute the determinant of the sub-matrix."""`

```
import numpy as np
# (a) Create a 4x4 NumPy array with numbers 1 to 16
arr = np.arange(1, 17).reshape(4, 4)
print("4x4 Array:\n", arr)

# (b) Extract the 2x2 bottom-left sub-matrix
sub_matrix = arr[2:, :2]
print("2x2 Bottom-left Sub-matrix:\n", sub_matrix)

# (c) Compute the determinant of the sub-matrix
det = np.linalg.det(sub_matrix)
print("Determinant of sub-matrix:", det)
```

4x4 Array:

```
[[ 1  2  3  4]
 [ 5  6  7  8]
 [ 9 10 11 12]
 [13 14 15 16]]
```

2x2 Bottom-left Sub-matrix:

```
[[ 9 10]
 [13 14]]
```

Determinant of sub-matrix: -4.0000000000000009

In [121...

```
"""Q3.(a) Create a DataFrame with 5 students and marks in 3 subjects.
(b) Add a column for total and average marks.
(c) Identify the topper and print their name with average."""
```

```
import pandas as pd
#(a) Create a DataFrame with 5 students and marks in 3 subjects
data = {
    'Name': ['Arpit', 'Mohan', 'Gagan', 'Chetan', 'Isha'],
    'Maths': [78, 85, 92, 88, 76],
    'Science': [82, 79, 95, 91, 80],
    'English': [74, 88, 90, 85, 78]
}
df = pd.DataFrame(data)
print("Student Marks DataFrame:\n", df)

#(b) Add total and average columns
df['Total'] = df[['Maths', 'Science', 'English']].sum(axis=1)
df['Average'] = df['Total'] / 3
print("\nDataFrame with Total & Average:\n", df)

#(c) Identify the topper and print their name with average
topper = df.loc[df['Average'].idxmax()]
print(f"\nTopper: {topper['Name']} | Average Marks: {topper['Average']:.2f}")
```

Student Marks DataFrame:

	Name	Maths	Science	English
0	Arpit	78	82	74
1	Mohan	85	79	88
2	Gagan	92	95	90
3	Chetan	88	91	85
4	Isha	76	80	78

DataFrame with Total & Average:

	Name	Maths	Science	English	Total	Average
0	Arpit	78	82	74	234	78.000000
1	Mohan	85	79	88	252	84.000000
2	Gagan	92	95	90	277	92.333333
3	Chetan	88	91	85	264	88.000000
4	Isha	76	80	78	234	78.000000

Topper: Gagan | Average Marks: 92.33

In [122...

```
"""Q4.(a) Simulate 1000 coin tosses using NumPy (1=Head, 0=Tail).
(b) Count frequency of heads and tails.
```

(c) Estimate probability of heads. Is it close to 0.5? Why/Why not?"""

```
import numpy as np
#(a) Simulate 1000 coin tosses (1=Head, 0=Tail)
tosses = np.random.randint(0, 2, 1000)
print("Coin Toss Results (first 25 shown):", tosses[:25])

#(b) Count frequency of heads and tails
heads = np.count_nonzero(tosses == 1)
tails = np.count_nonzero(tosses == 0)
print(f"Heads: {heads}, Tails: {tails}")

#(c) Estimate probability of heads
prob_heads = heads / len(tosses)
print(f"Probability of Heads: {prob_heads:.3f}")

#Commenting on result
if abs(prob_heads - 0.5) < 0.05:
    print("Yes, it's close to 0.5 – randomness averages out over many tosses.")
else:
    print("Not very close to 0.5 – randomness can still cause small deviations")
```

Coin Toss Results (first 25 shown): [1 0 0 0 0 1 1 0 1 0 1 1 0 1 0 0 0 0 0 1 1 0 0 0 0]

Heads: 505, Tails: 495

Probability of Heads: 0.505

Yes, it's close to 0.5 – randomness averages out over many tosses.

In [123...

"""Q5.(a) Create a DataFrame of employees with columns: ID, Name, Salary.  
(b) Add a Bonus column = 10% of Salary.  
(c) Display employees with salary above average."""

```
#(a) Create a DataFrame of employees with columns: ID, Name, Salary.
import pandas as pd
data = {
    'ID': [101, 102, 103, 104, 105],
    'Name': ['Arpit', 'Gagan', 'Mohan', 'Chetan', 'Isha'],
    'Salary': [65000, 52000, 48000, 60000, 55000]
}
df = pd.DataFrame(data)
print("Employees DataFrame:\n",df)

# (b) Add Bonus column = 10% of Salary
df['Bonus'] = df['Salary'] * 0.10
print("\nDataFrame with Bonus:\n", df)

# (c) Display employees with salary above average
avg_salary = df['Salary'].mean()
above_avg = df[df['Salary'] > avg_salary]
print(f"\nAverage Salary: {avg_salary:.2f}")
print("\nEmployees with Salary above Average:\n", above_avg)
```

Employees DataFrame:

	ID	Name	Salary
0	101	Arpit	65000
1	102	Gagan	52000
2	103	Mohan	48000
3	104	Chetan	60000
4	105	Isha	55000

DataFrame with Bonus:

	ID	Name	Salary	Bonus
0	101	Arpit	65000	6500.0
1	102	Gagan	52000	5200.0
2	103	Mohan	48000	4800.0
3	104	Chetan	60000	6000.0
4	105	Isha	55000	5500.0

Average Salary: 56000.00

Employees with Salary above Average:

	ID	Name	Salary	Bonus
0	101	Arpit	65000	6500.0
3	104	Chetan	60000	6000.0

In [124...

```
"""Q6.(a) Create a 3x3 NumPy array with values 1 to 9.
(b) Find its transpose and inverse.
(c) Verify that  $A \times A^{-1} \approx I$ ."""

import numpy as np

# (a) Create a 3x3 NumPy array with values 1 to 9
A = np.arange(1, 10).reshape(3, 3)
print("Matrix A:\n", A)

# (b) Find its transpose and inverse
A_T = A.T
try:
    A_inv = np.linalg.inv(A)
    print("\nTranspose of A:\n", A_T)
    print("\nInverse of A:\n", A_inv)
except np.linalg.LinAlgError:
    print("\nMatrix A is singular and cannot be inverted.")

# (c) Verify that  $A \times A^{-1} \approx I$ 
if 'A_inv' in locals():
    I_approx = np.dot(A, A_inv)
    print("\n $A \times A^{-1} \approx I$ :\n", I_approx)
```

Matrix A:

```
[[1 2 3]
 [4 5 6]
 [7 8 9]]
```

Matrix A is singular and cannot be inverted.

In [125... `"""Q7.(a) Generate random daily temperatures (30 values, range 20–40°C).  
(b) Find the hottest and coldest day.  
(c) Compute mean, median, and standard deviation of temperatures."""`

```
import numpy as np

# (a) Generate random daily temperatures (30 values, range 20–40°C)
temps = np.random.uniform(20, 40, 30)
print("Daily Temperatures (°C):\n", np.round(temps, 2))

# (b) Find the hottest and coldest day
hottest = np.max(temps)
coldest = np.min(temps)
print(f"\nHottest Temperature: {hottest:.2f}°C")
print(f"Coldest Temperature: {coldest:.2f}°C")

# (c) Compute mean, median, and standard deviation
mean_temp = np.mean(temps)
median_temp = np.median(temps)
std_temp = np.std(temps)
print(f"\nMean: {mean_temp:.2f}°C")
print(f"Median: {median_temp:.2f}°C")
print(f"Standard Deviation: {std_temp:.2f}")
```

Daily Temperatures (°C):  
[38. 28.09 37.91 21.3 33.71 24.62 30.78 35.46 34.93 34.43 29.66 23.88  
24.11 30.71 28.76 39.5 26.02 21.58 34.54 39.51 35.64 24.94 20.73 28.14  
37.56 28.64 25.68 38.71 35.45 32.22]

Hottest Temperature: 39.51°C  
Coldest Temperature: 20.73°C

Mean: 30.84°C  
Median: 30.74°C  
Standard Deviation: 5.77

In [126... `"""Q8.(a) Create a Pandas Series of marks for 8 students.  
(b) Replace all marks below 40 with "Fail".  
(c) Count how many students passed."""`

```
import pandas as pd

# (a) Create a Pandas Series of marks for 8 students
marks = pd.Series([55, 38, 72, 91, 47, 30, 65, 84])
print("Marks Series:\n", marks)

# (b) Replace all marks below 40 with "Fail"
marks_updated = marks.apply(lambda x: "Fail" if x < 40 else x)
print("\nUpdated Marks (with 'Fail'):\n", marks_updated)

# (c) Count how many students passed
passed_count = sum(marks_updated != "Fail")
print(f"\nNumber of Students Passed: {passed_count}")
```

Marks Series:

```
0    55
1    38
2    72
3    91
4    47
5    30
6    65
7    84
```

dtype: int64

Updated Marks (with 'Fail'):

```
0    55
1    Fail
2    72
3    91
4    47
5    Fail
6    65
7    84
```

dtype: object

Number of Students Passed: 6

In [127...

```
"""Q9.(a) Generate 500 random integers between 1 and 6 (simulate dice rolls).
(b) Count how many times each face appears.
(c) Compute relative frequencies and compare with theoretical 1/6."""
```

```
import numpy as np
```

```
# (a) Generate 500 random integers between 1 and 6 (simulate dice rolls)
rolls = np.random.randint(1, 7, 500)
print("First 20 Dice Rolls:", rolls[:20])
```

```
# (b) Count how many times each face appears
counts = {face: np.count_nonzero(rolls == face) for face in range(1, 7)}
print("\nFace Counts:", counts)
```

```
# (c) Compute relative frequencies and compare with theoretical 1/6
rel_freq = {face: count / 500 for face, count in counts.items()}
print("\nRelative Frequencies:", rel_freq)
print("\nTheoretical Probability (1/6) ≈ 0.1667")
for face, freq in rel_freq.items():
    print(f"Face {face}: {freq:.3f} (diff = {abs(freq - 1/6):.3f})")
```

First 20 Dice Rolls: [3 6 5 4 5 2 1 1 2 1 4 6 3 1 2 3 2 4 2 6]

Face Counts: {1: 93, 2: 74, 3: 81, 4: 81, 5: 86, 6: 85}

Relative Frequencies: {1: 0.186, 2: 0.148, 3: 0.162, 4: 0.162, 5: 0.172, 6: 0.17}

Theoretical Probability ( $1/6$ )  $\approx$  0.1667

Face 1: 0.186 (diff = 0.019)

Face 2: 0.148 (diff = 0.019)

Face 3: 0.162 (diff = 0.005)

Face 4: 0.162 (diff = 0.005)

Face 5: 0.172 (diff = 0.005)

Face 6: 0.170 (diff = 0.003)

```
In [128... """Q10.(a) Create a DataFrame with 6 products (Name, Quantity, Price).
(b) Add a column "Total = Quantity × Price".
(c) Find which product generated maximum sales revenue."""

import pandas as pd

# (a) Create a DataFrame with 6 products
data = {
    'Name': ['Pen', 'Notebook', 'Eraser', 'Marker', 'Stapler', 'Folder'],
    'Quantity': [50, 30, 100, 20, 15, 40],
    'Price': [5, 20, 2, 15, 50, 10]
}
df = pd.DataFrame(data)
print("Products DataFrame:\n", df)

# (b) Add a column "Total = Quantity × Price"
df['Total'] = df['Quantity'] * df['Price']
print("\nDataFrame with Total Sales:\n", df)

# (c) Find which product generated maximum sales revenue
max_sales_product = df.loc[df['Total'].idxmax()]
print(f"\nProduct with Maximum Sales: {max_sales_product['Name']} | Revenue: {
```

Products DataFrame:

	Name	Quantity	Price
0	Pen	50	5
1	Notebook	30	20
2	Eraser	100	2
3	Marker	20	15
4	Stapler	15	50
5	Folder	40	10

DataFrame with Total Sales:

	Name	Quantity	Price	Total
0	Pen	50	5	250
1	Notebook	30	20	600
2	Eraser	100	2	200
3	Marker	20	15	300
4	Stapler	15	50	750
5	Folder	40	10	400

Product with Maximum Sales: Stapler | Revenue: 750

```
In [129... """Q11.(a) Create a DataFrame with some missing values (NaN).
(b) Fill missing values with column mean.
(c) Drop rows where more than 1 value is missing."""

import pandas as pd
import numpy as np

# (a) Create a DataFrame with some missing values
data = {
    'A': [10, 20, np.nan, 40, np.nan],
    'B': [5, np.nan, 15, np.nan, 25],
    'C': [11, np.nan, 12, 16, 20]
}
df = pd.DataFrame(data)
print("Original DataFrame:\n", df)

# (b) Fill missing values with column mean
df_filled = df.fillna(df.mean())
print("\nDataFrame after filling NaNs with column mean:\n", df_filled)

# (c) Drop rows where more than 1 value is missing
df_dropped = df.dropna(thresh=df.shape[1]-1) # keep rows with at least 2 non-
print("\nDataFrame after dropping rows with >1 missing value:\n", df_dropped)
```



Original DataFrame:

	A	B	C
0	10.0	5.0	11.0
1	20.0	NaN	NaN
2	NaN	15.0	12.0
3	40.0	NaN	16.0
4	NaN	25.0	20.0

DataFrame after filling NaNs with column mean:

	A	B	C
0	10.000000	5.0	11.00
1	20.000000	15.0	14.75
2	23.333333	15.0	12.00
3	40.000000	15.0	16.00
4	23.333333	25.0	20.00

DataFrame after dropping rows with >1 missing value:

	A	B	C
0	10.0	5.0	11.0
2	NaN	15.0	12.0
3	40.0	NaN	16.0
4	NaN	25.0	20.0

In [130...

```
"""Q12.(a) Create a NumPy array with integers from 1 to 30.  
(b) Reshape it into a 5x6 matrix.  
(c) Extract all even numbers from the matrix and compute their average."""
```

```
import numpy as np
```

```
# (a) Create a NumPy array with integers from 1 to 30  
arr = np.arange(1, 31)  
print("Array:", arr)
```

```
# (b) Reshape into a 5x6 matrix  
matrix = arr.reshape(5, 6)  
print("\n5x6 Matrix:\n", matrix)
```

```
# (c) Extract all even numbers and compute their average  
evens = matrix[matrix % 2 == 0]  
average_evens = np.mean(evens)  
print("\nEven Numbers:", evens)  
print("Average of Even Numbers:", average_evens)
```

```
Array: [ 1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24
       25 26 27 28 29 30]
```

5x6 Matrix:

```
[[ 1  2  3  4  5  6]
 [ 7  8  9 10 11 12]
 [13 14 15 16 17 18]
 [19 20 21 22 23 24]
 [25 26 27 28 29 30]]
```

Even Numbers: [ 2 4 6 8 10 12 14 16 18 20 22 24 26 28 30]

Average of Even Numbers: 16.0

In [131...

```
"""Q13.
(a) Create a DataFrame of 6 students with columns: Name, Age, Marks.
(b) Select only students who scored above the overall average marks.
(c) Display names of students younger than 20 whose marks are above 60"""

import pandas as pd

# (a) Create a DataFrame of 6 students
data = {
    'Name': ['Arpit', 'Bhaskar', 'Gagan', 'Mohan', 'Isha', 'Farhan'],
    'Age': [19, 21, 18, 22, 20, 17],
    'Marks': [65, 58, 72, 49, 80, 68]
}
df = pd.DataFrame(data)
print("Students DataFrame:\n", df)

# (b) Select students who scored above overall average marks
avg_marks = df['Marks'].mean()
above_avg = df[df['Marks'] > avg_marks]
print(f"\nOverall Average Marks: {avg_marks:.2f}")
print("Students with Marks above Average:\n", above_avg)

# (c) Display names of students younger than 20 whose marks are above 60
young_high = df[(df['Age'] < 20) & (df['Marks'] > 60)]
print("\nStudents younger than 20 with Marks > 60:\n", young_high['Name'].tolist())
```

Students DataFrame:

	Name	Age	Marks
0	Arpit	19	65
1	Bhaskar	21	58
2	Gagan	18	72
3	Mohan	22	49
4	Isha	20	80
5	Farhan	17	68

Overall Average Marks: 65.33

Students with Marks above Average:

	Name	Age	Marks
2	Gagan	18	72
4	Isha	20	80
5	Farhan	17	68

Students younger than 20 with Marks > 60:

['Arpit', 'Gagan', 'Farhan']

In [132...

```
"""Q14.(a) Create a Pandas DataFrame with 5 employees having columns: Name, Ta
Hours_Worked.
(b) Add a new column Efficiency = Tasks_Completed / Hours_Worked.
(c) Identify the employee with the highest efficiency and print their name and

import pandas as pd

# (a) Create a DataFrame with 5 employees
data = {
    'Name': ['Arpit', 'Mohan', 'Chetan', 'Bhasker', 'Isha'],
    'Tasks_Completed': [50, 40, 60, 55, 45],
    'Hours_Worked': [10, 8, 15, 12, 9]
}
df = pd.DataFrame(data)
print("Employees DataFrame:\n", df)

# (b) Add a new column Efficiency = Tasks_Completed / Hours_Worked
df['Efficiency'] = df['Tasks_Completed'] / df['Hours_Worked']
print("\nDataFrame with Efficiency:\n", df)

# (c) Identify the employee with the highest efficiency
top_employee = df.loc[df['Efficiency'].idxmax()]
print(f"\nHighest Efficiency: {top_employee['Name']} | Value: {top_employee['E
```

Employees DataFrame:

	Name	Tasks_Completed	Hours_Worked
0	Arpit	50	10
1	Mohan	40	8
2	Chetan	60	15
3	Bhasker	55	12
4	Isha	45	9

DataFrame with Efficiency:

	Name	Tasks_Completed	Hours_Worked	Efficiency
0	Arpit	50	10	5.000000
1	Mohan	40	8	5.000000
2	Chetan	60	15	4.000000
3	Bhasker	55	12	4.583333
4	Isha	45	9	5.000000

Highest Efficiency: Arpit | Value: 5.00

## Part C - Case Study

In [133...

```
"""Case Study 1 – Student Performance Analytics
You are given data of 100 students containing their marks in Math, Science, and English.
(a) Using Pandas, calculate the average marks for each student and identify the top 5 performers.
(b) Compute subject-wise average and standard deviation; comment which subject has the highest variation.
(c) Draw a boxplot for each subject and comment on outliers."""

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

# Simulating the data for 100 students
np.random.seed(0) # for reproducibility
data = {
    'Student': [f'Student_{i+1}' for i in range(100)],
    'Math': np.random.randint(40, 101, 100),
    'Science': np.random.randint(35, 101, 100),
    'English': np.random.randint(30, 101, 100)
}
df = pd.DataFrame(data)

# (a) Calculate average marks for each student and identify top 5 performers
df['Average'] = df[['Math', 'Science', 'English']].mean(axis=1)
top5 = df.nlargest(5, 'Average')
print("Top 5 Students based on Average Marks:\n", top5[['Student', 'Average']])

# (b) Compute subject-wise average and standard deviation
subject_stats = df[['Math', 'Science', 'English']].agg(['mean', 'std'])
print("\nSubject-wise Average and Standard Deviation:\n", subject_stats)
highest_variation_subject = subject_stats.loc['std'].idxmax()
print(f"\nSubject with highest variation: {highest_variation_subject}")
```

```
# (c) Draw a boxplot for each subject
plt.figure(figsize=(8,5))
df[['Math', 'Science', 'English']].boxplot()
plt.title("Boxplot of Student Marks")
plt.ylabel("Marks")
plt.show()

print("\nComment on outliers: Points outside the whiskers in the boxplot are c
```

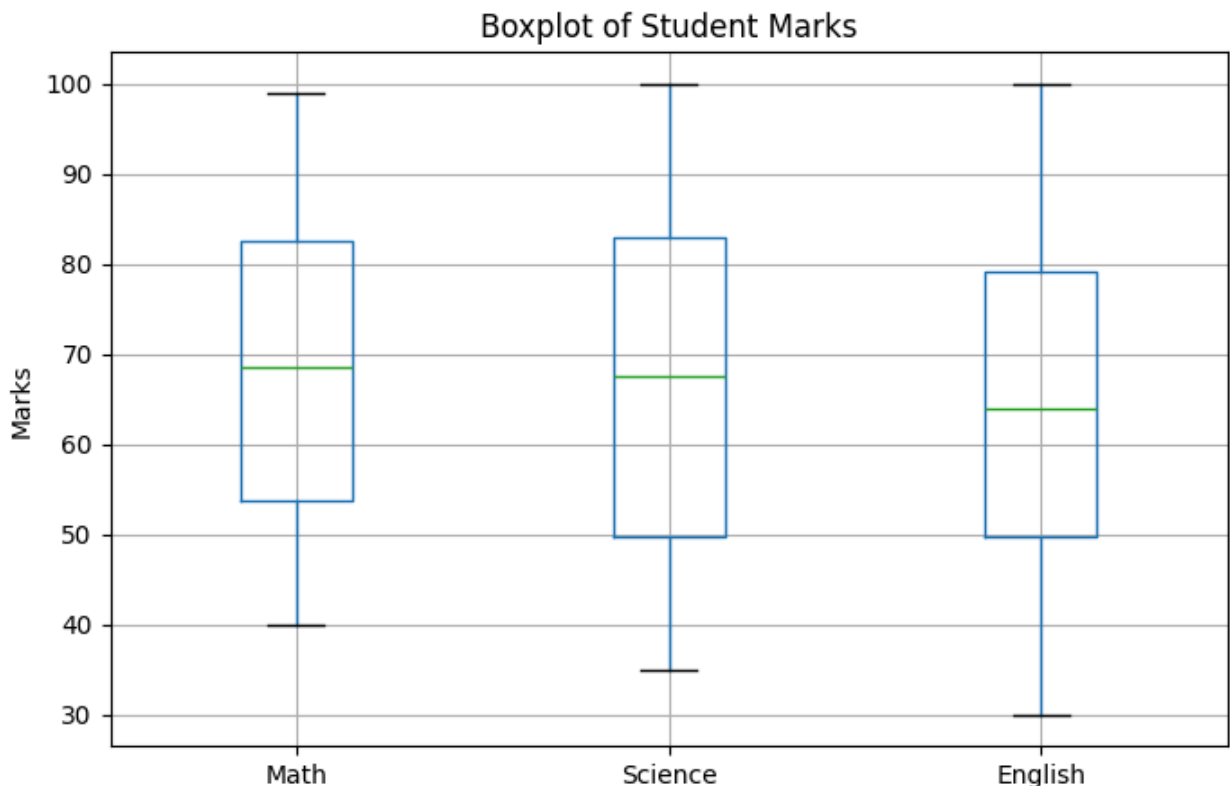
Top 5 Students based on Average Marks:

	Student	Average
97	Student_98	92.333333
84	Student_85	91.333333
53	Student_54	88.666667
77	Student_78	87.000000
67	Student_68	86.666667

Subject-wise Average and Standard Deviation:

	Math	Science	English
mean	68.620000	66.860000	64.410000
std	18.117897	19.115819	19.664918

Subject with highest variation: English



Comment on outliers: Points outside the whiskers in the boxplot are outliers, indicating exceptionally low or high marks in that subject.

In [134...

```
"""Case Study 2 – Sales Data Exploration
A company records sales transactions in a dataset with columns: Product, Quantity, Price, Region.
(a) Add a new column Total = Quantity × Price and compute overall revenue.
```

- (b) Group data by region and identify which region contributes the most sales.  
(c) Plot a histogram of total sales per product and discuss which products are

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

#Simulating sales data
np.random.seed(0) # for reproducibility
products = ['Product_A', 'Product_B', 'Product_C', 'Product_D', 'Product_E']
regions = ['North', 'South', 'East', 'West']
data = {
    'Product': np.random.choice(products, 100),
    'Quantity': np.random.randint(1, 21, 100),
    'Price': np.random.randint(50, 501, 100),
    'Region': np.random.choice(regions, 100)
}
df = pd.DataFrame(data)

#(a)Add Total column and compute overall revenue
df['Total'] = df['Quantity'] * df['Price']
overall_revenue = df['Total'].sum()
print("Overall Revenue:", overall_revenue)

#(b)Group by region and identify top contributing region
region_sales = df.groupby('Region')['Total'].sum()
top_region = region_sales.idxmax()
print("\nSales by Region:\n", region_sales)
print(f"Region with highest sales: {top_region}")

#(c)Plot histogram of total sales per product
product_sales = df.groupby('Product')['Total'].sum()
plt.figure(figsize=(8,5))
product_sales.plot(kind='bar', color='skyblue')
plt.title("Total Sales per Product")
plt.ylabel("Total Sales")
plt.show()


print("\nComment on underperforming products: Products with significantly lowe
```

Overall Revenue: 303042

Sales by Region:

Region	Total
East	90107
North	90538
South	48222
West	74175

Name: Total, dtype: int64  
Region with highest sales: North



Comment on underperforming products: Products with significantly lower total sales compared to others may be underperforming.

In [135...

```
"""Case Study 3 – Predicting Fuel Efficiency (Auto MPG Dataset)
You are analyzing the Auto MPG dataset (UCI Repository) with features: mpg, cy
displacement, horsepower, weight, acceleration, model year, origin.
(a) Compute the correlation of mpg with each numeric feature and discuss which
influence fuel efficiency most.
(b) Plot a scatter plot of weight vs mpg; describe the trend you observe."""

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

#(a) Computing the correlation of mpg with each numeric feature and discuss wh
# Loading the Auto MPG dataset

url = 'https://archive.ics.uci.edu/ml/machine-learning-databases/auto-mpg/auto
column_names = ['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
                 'acceleration', 'model_year', 'origin']

# Use raw string for regex separator and handle missing values
df = pd.read_csv(url, sep=r'\s+', names=column_names)
# Convert 'horsepower' to numeric, coercing errors to NaN
df['horsepower'] = pd.to_numeric(df['horsepower'], errors='coerce')
```

```

# Drop rows with NaN values in 'horsepower' after coercion
df.dropna(subset=['horsepower'], inplace=True)

# Now correlation will work
corr_matrix = df.corr(numeric_only=True) # Calculate correlation only on numeric
print("Correlation with MPG:\n", corr_matrix['mpg'].sort_values(ascending=False))

#(b) Scatter Plot of Weight vs MPG
# Scatter plot of weight vs mpg
plt.figure(figsize=(8, 6))
sns.scatterplot(x='weight', y='mpg', data=df, color='purple')
plt.title('Scatter Plot: Weight vs MPG')
plt.xlabel('Weight (lbs)')
plt.ylabel('Miles per Gallon (MPG)')
plt.show()

```

Correlation with MPG:

mpg	1.000000
cylinders	0.950721
horsepower	0.896017
acceleration	-0.348746
weight	-0.505419
model_year	-0.562543

Name: mpg, dtype: float64



Scatter Plot: Weight vs MPG

